

# CS-2006

# Operating System

## Lecture 6

Mahzaib Younas

Lecturer Department of Computer Science

FAST NUCES CFD

# Outlines

- Overview
- Multicore Programming
- Multithreading Models
- Thread Libraries
- Threading Issues

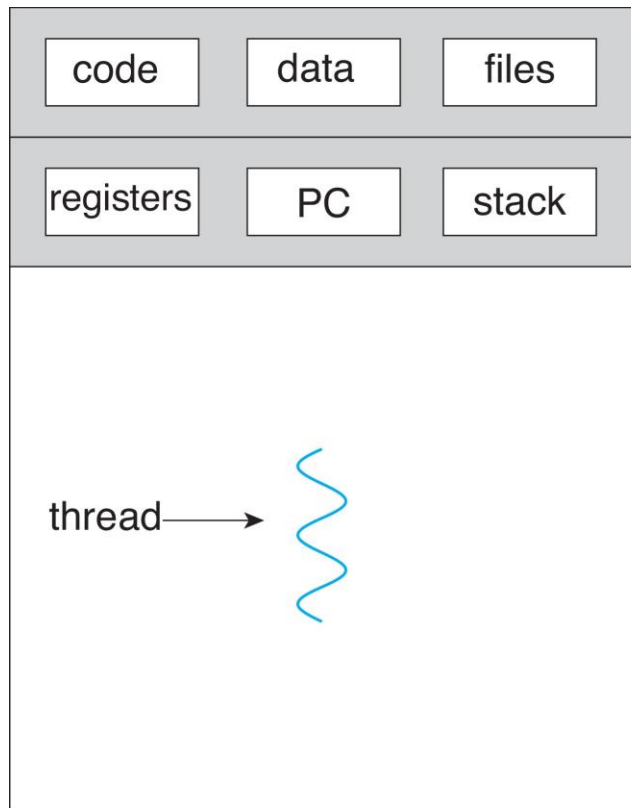
# Overview:

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
  - Update display
  - Fetch data
  - Spell checking
  - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

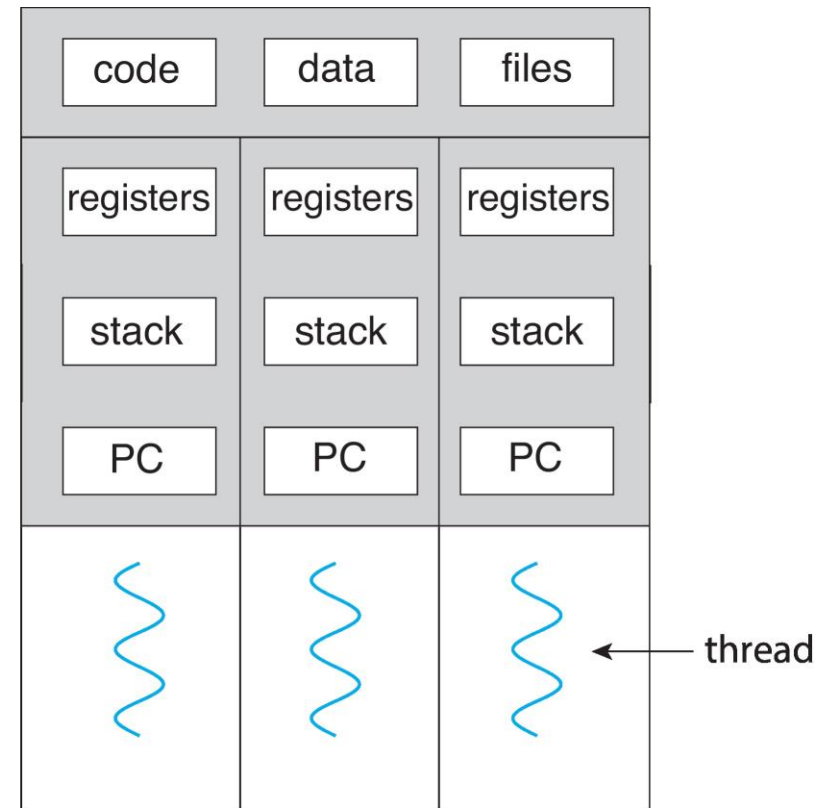
# Threading

- A thread is a basic unit of CPU utilization; it comprises
  - thread ID
  - program counter (PC)
  - register set
  - stack.
- It shares with other threads belonging to the same process its
  - code section
  - data section
  - operating-system resources
- For Example: open files and signals.
- A traditional process has a single thread of control. If a process has multiple threads of control, it can perform more than one task at a time.

# Threading



single-threaded process

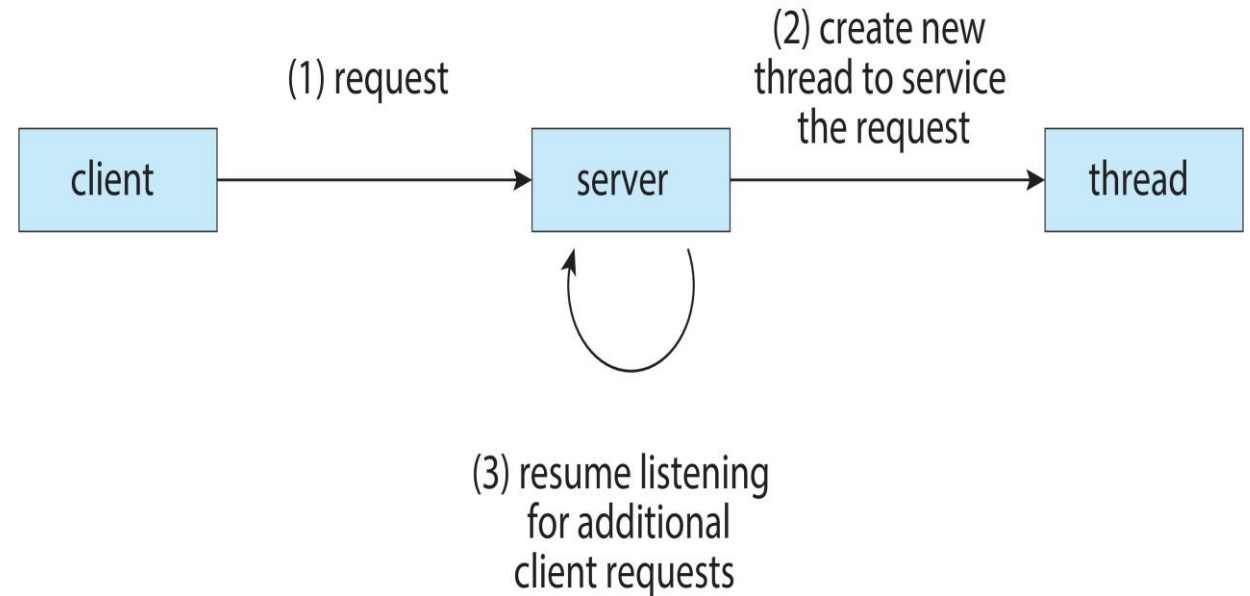


multithreaded process

# Multi Threading Server Architecture

## *Example*

If the web-server process is multithreaded the server will create a separate thread that listens for client requests. When a request is made, rather than creating another process, the server creates a new thread to service the request and resumes listening for additional requests



# Benefits of Multi Threading

- **Responsiveness**

- may allow **continued execution** if part of process is blocked, especially important for user interfaces.
- Mouse Click Response

- **Resource Sharing**

- threads share resources of process, **easier than shared memory or message passing**

- **Economy**

- **cheaper than process creation**, thread switching **lower overhead than context switching**

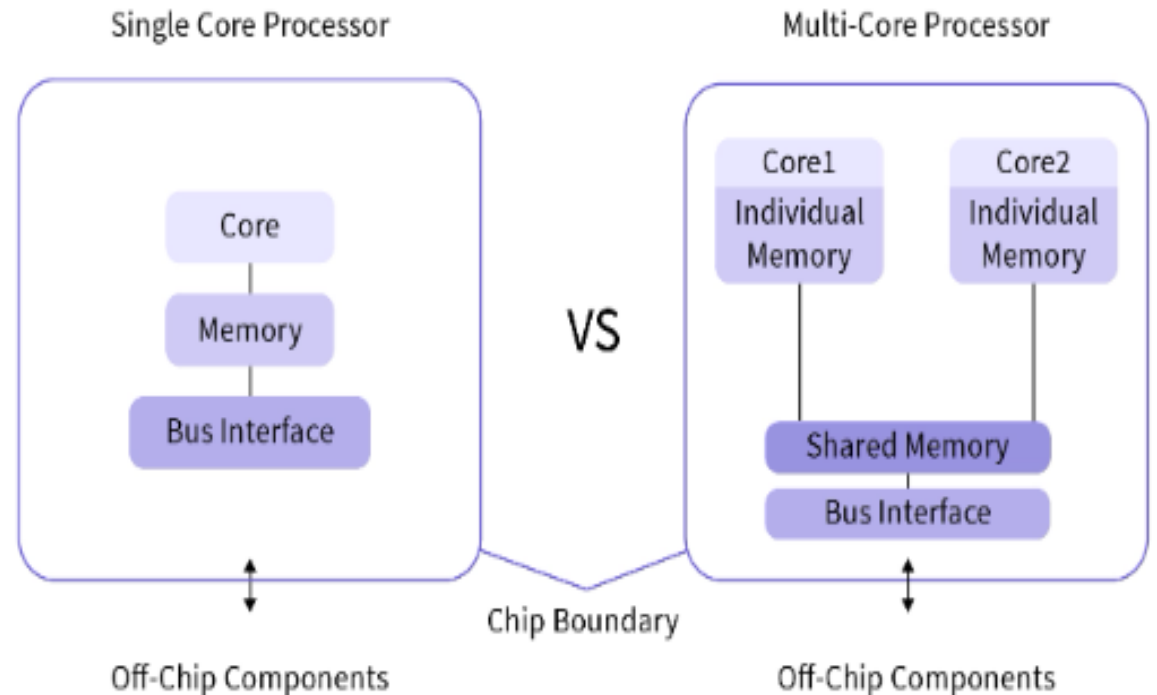
- **Scalability**

- process can take advantage of **multicore architectures**
- A single-threaded process can run on only one processor, regardless how many are available

# Multi Core Programming

A multicore processor system is a single processor with multiple execution cores in one chip.

- Multicore or multiprocessor systems puts pressure on programmers, challenges include:
  - **Dividing activities**
  - **Balance**
  - **Data splitting**
  - **Data dependency**
  - **Testing and debugging**





# Multicore Programming Challenges

- **Identifying Task / Dividing Activities**

Examine the application to find area that can be divided into separate, concurrent task. Ideally, tasks are independent of one another and this can run in parallel cores.

- **Balance**

- Identifying tasks that can run in parallel, programmers must also ensure that the tasks perform equal work of equal value.

- **Data Splitting**

- Applications are divided into separate tasks, the data accessed and manipulated by the tasks must be divided to run on separate cores

# Multicore Programming Challenges

- **Data Dependency**

The data accessed by the tasks must be examined for dependencies between two or more tasks. When **one task depends on data from another**, **programmers must ensure that the execution of the tasks is synchronized** to accommodate the data dependency


- **Testing and debugging**

A program is running **in parallel on multiple cores**, many different execution paths are possible. **Testing and debugging such concurrent programs is inherently more difficult than testing and debugging single-threaded applications**

# Multi Core Programming (Cont..)


## Parallelism

- implies a system can perform **more than one task simultaneously**
- Execute on multi core system



## Concurrency


- supports more than one task making progress **Single processor / core, scheduler providing concurrency**
- Execute on single core System



# Multi Core Programming (Cont...)


## Data Parallelism

distributes subsets of the same data  
across multiple cores, same operation  
on each



## Task Parallelism

distributing threads across cores, each  
thread performing unique operation



# Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- $S$  is serial portion
- $N$  processing cores


$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As  $N$  approaches infinity, speedup approaches  $1 / S$

**Serial portion of an application has disproportionate effect on performance gained by adding additional cores**

- But does the law take into account contemporary multicore systems?

# Amdahl Law's (Cont...)



# User Threads and Kernel Threads

- **User threads**

- management done by user-level threads library

- Three primary thread libraries:


- POSIX **Pthreads**
  - Windows threads
  - Java threads

- **Kernel Thread:**

Supported by the Kernel

- Examples – virtually all general-purpose operating systems, including:
    - Windows
    - Linux
    - Mac OS X
    - iOS
    - Android

# User Threads and Kernel Threads (cont...)





# Threading Models


A relationship **must exist between user threads and kernel threads**

- Three common ways of establishing such a relationship:
  - many-to-one model
  - one-to one model
  - many-to-many model

# Many to One Model:

Many user-level threads mapped to single kernel thread


- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
  - **Solaris Green Threads**
  - **GNU Portable Threads**



# One to One Model


Each user-level thread maps to kernel thread

- Creating a **user-level thread creates a kernel thread**
- **More concurrency** than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
  - Windows
  - Linux




# Many to Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the *ThreadFiber* package
- Otherwise not very common



# Two Mode Model

- Similar to Many to Many
- except that it allows a user thread to be **bound** to kernel thread



# Thread Libraries

- Thread library provides **programmer with API for creating and managing threads**
- Two primary ways of implementing
  - Library entirely in user space
  - Kernel-level library supported by the OS

# Types of Libraries

- Three main libraries are used
  - POSIX Pthreads
    - The threads extension of the POSIX standard, may be provided as either a user-level or a kernel-level library
  - Windows
    - Windows thread library is a kernel-level library available on Windows systems.
  - Java
    - Java thread API allows threads to be created and managed directly in Java programs.

# Pthreads

- A POSIX **standard (IEEE 1003.1c)** API for thread creation and **synchronization**
- Specification, not implementation
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)



# How to create a Thread in C

- **pthread\_create**

- It is used to create a new thread.

- Parameters Detail

- Thread

Pointer to an unsigned integer value that return the **thread ID** of the thread that is created.

- Attr

Pointer to a structure that is used to define the **thread attributes** like state, scheduling policy

- Start-Routine

Pointer to subroutine that is **executed** by thread. It return types must be void.

- Args

Pointer to void that contain the arguments to the function which defined earlier

```
int pthread_create(pthread_t * thread,  
                  const pthread_attr_t * attr,  
                  void * (*start_routine)(void *),  
                  void *arg);
```

# Pthreads Code Example:

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```

# Pthread Example

```
/* The thread will execute in this function */  
void *runner(void *param)  
{  
    int i, upper = atoi(param);  
    sum = 0;  
  
    for (i = 1; i <= upper; i++)  
        sum += i;  
  
    pthread_exit(0);  
}
```

# Function of Threads

Function Name	Purpose
Pthread_Create	Used to <b>Create</b> the thread
Pthread_exit	Used to <b>terminate</b> the thread
Pthread_join	Used to <b>wait for termination</b> of thread
Pthread_self	Used to <b>get the id</b> of current running thread
Pthread_cancel	Used to send <b>the cancellation request</b>

# Practice Code

Write a program in C by using threading, in which threads prints 0-4 while main process prints 20-24.

# Code

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
```

## //Prototyping of thread

```
Void *thread_print( void *arg)
int i, j;
```

```
int main ( )
{
    Pthread_t a; //Thread declaration
    Pthread_create(&a, NULL, thread_print, NULL);
    Pthread_join(a, NULL);
    Printf("print 4 no of given range");
    For(j=20;j<25;j++)
    {
        Printf("%d\n", j);
    }

    Void *thread_print( void *arg)
    Printf("inside thread");
    For(i=1; i<5; i++)
    {
        Printf("%d\n", i);
    }
}
```

# Practice Code

Write a program that creates threads based on the input given by user. Each thread should execute function print () and display its thread ID. The output should be like:

Hello I am thread 1 my ID is 123

Hello I am thread 2 my ID is 234....

The main thread should wait for the child threads to terminate and then call exit.

# Code

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
void *print(void *threadid)
{
    pthread_t tid = pthread_self();
    int id = *((int *)threadid);

    printf("Hello -> I AM THREAD %d AND MY ID IS  
= %d\n", id, tid);

    pthread_exit(NULL);
}
```

```
int main()
{
    int num;
    printf("ENTER THE NUMBER OF THREADS TO CREATE = ");
    scanf("%d", &num);

    pthread_t threads[num];
    int threadids[num];
    for (int i = 0; i < num; i++)
    {
        threadids[i] = i + 1;
        pthread_create(&threads[i], NULL, print, &threadids[i]);
    }
    for (int i = 0; i < num; i++)
    {
        pthread_join(threads[i], NULL);
    }
    printf("ALL THREADS HAVE BEEN EXECUTED  
SUCCESSFULLY...\n");
    exit(0);
}
```



# Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal handling
  - Synchronous and asynchronous
- Thread cancellation of target thread
  - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations

# Signal Handling (Cont.)

- Where should a signal be delivered for multi-threaded?
  - Deliver the signal to the thread to which the signal applies
  - Deliver the signal to every thread in the process
  - Deliver the signal to certain threads in the process
  - Assign a specific thread to receive all signals for the process

# Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
  - **Asynchronous cancellation** terminates the **target thread immediately**
  - **Deferred cancellation** allows the target thread to **periodically check if it should be cancelled**
- Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```

# Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but **actual cancellation depends on thread state**

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous


- If thread has **cancellation disabled, cancellation remains pending until** thread enables it
- Default type is **deferred**
  - Cancellation only occurs when thread reaches **cancellation point**
    - i.e., `pthread_testcancel()`
    - Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals

# Thread-Local Storage

- **Thread-local storage (TLS)** allows each thread to have its own copy of data
- Useful when you do not have control over the **thread creation process** (i.e., **when using a thread pool**)
- Different from local variables
  - Local variables visible **only during single function invocation**
  - TLS visible across function invocations
- Similar to **static** data
  - **TLS is unique to each thread**

# Scheduler Activations

- Both **M:M and Two-level models** require communication to maintain the appropriate number of **kernel threads allocated to the application**
- Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
  - Appears to be a **virtual processor on which process can schedule user thread to run**
  - Each LWP attached to kernel thread
  - How many LWPs to create?
- Scheduler activations provide **upcalls - a communication mechanism from the kernel to the upcall handler in the thread library**
- This communication allows an application to maintain the correct number kernel threads



# Linux Threads

- Linux refers to them as *tasks* rather than *threads*
- Thread creation is done through **clone()** system call
- clone() allows a child task to share the address space of the parent task (process)
  - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- struct task\_struct points to process data structures (shared or unique)