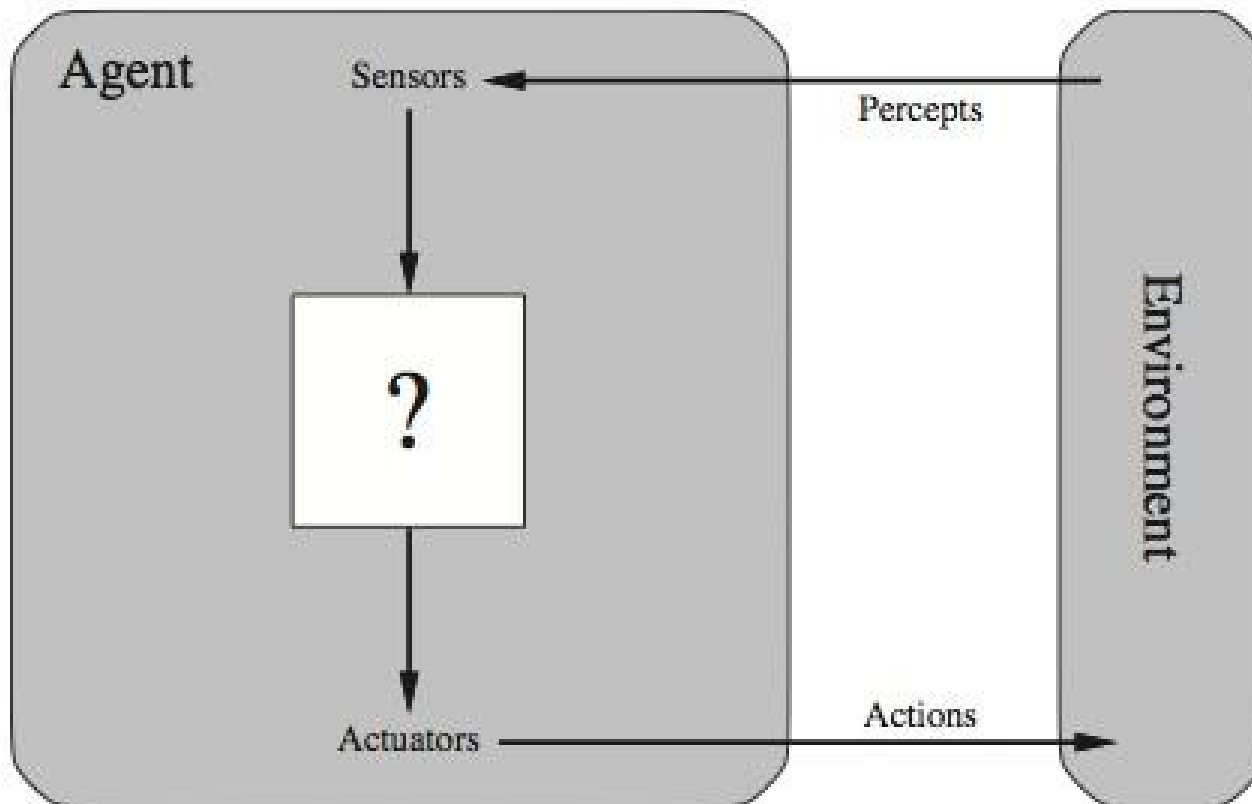


AI 2002

Artificial Intelligence

Dr. Hashim Yasin

Agent with an Environment



Task Environments

- ▶ Environment are essentially the "**problems**" to which rational agents are the "**solutions**".

Rationality

What is rational at any given time?

It depends on four things:

- ▶ The **performance measure** that defines the criterion of success.
- ▶ The **agent's prior knowledge** of the **environment**.
- ▶ The **actions** that the agent can perform.
- ▶ The agent's **percept sequence** to date.

Task Environments

- ▶ In designing an agent, the first step must always be to specify the task environment (PEAS) as fully as possible.

PEAS:

- ▶ Performance measure
- ▶ Environment
- ▶ Actuators
- ▶ Sensors

Task Environments ... Examples

PEAS for an automated taxi-driving

- ▶ **Performance measure:** Safe, fast, legal, comfortable trip, maximize profits, etc.
- ▶ **Environment:** Roads, other traffic, pedestrians, customers, etc.
- ▶ **Actuators:** Steering wheel, accelerator, brakes, signal, horn, etc.
- ▶ **Sensors:** Cameras, speedometer, GPS, odometer, engine sensors, keyboard, etc.

Task Environments ... Examples

PEAS for medical diagnosis system

- ▶ **Performance measure:** Healthy patient, minimize costs etc.
- ▶ **Environment:** Patient, hospital, staff etc.
- ▶ **Actuators:** Screen display (questions, tests, diagnoses, treatments, referrals)
- ▶ **Sensors:** Keyboard (entry of symptoms, findings, patient's answers)

Task Environments ... Examples

PEAS for satellite image analysis system

- ▶ **Performance measure:** correct image categorization
- ▶ **Environment:** downlink from the orbiting satellite
- ▶ **Actuators:** display categorization of scene
- ▶ **Sensors:** colour pixel arrays

Environment Types

Environment Types

- ▶ **Fully observable** vs. **partially observable**
- ▶ **Deterministic** vs. **stochastic**
- ▶ **Episodic** vs. **sequential**
- ▶ **Static** vs. **dynamic**
- ▶ **Discrete** vs. **continuous**
- ▶ **Single agent** vs. **multi-agent**

Environment Types

Fully observable vs. partially observable

Fully observable:

- If an agent's sensors give it access to the **complete** state of the environment at each point in time.
- A task environment is **effectively fully observable** if *the sensors detect all aspects that are relevant to the choice of action.*
- **Convenient**, because the agent need not to maintain any internal state to keep track of the world

Partially observable:

- **Parts** of the state are simply missing from sensor data
- **Noisy and inaccurate** sensors
 - A **vacuum agent** with only a local dirt sensor cannot tell whether there is dirt in other squares
 - An **automated taxi** cannot see what other drivers are thinking.

Environment Types

Deterministic vs. stochastic

Deterministic:

- If the next state of the environment is completely determined by the current state and the action executed by the agent.
- Vacuum-cleaner world is deterministic.

Stochastic:

- If the environment is partially observable then it could be stochastic.
- Taxi driving is clearly stochastic in this sense, because one can never predict the behaviour of the traffic exactly.

Environment Types

Episodic vs. Sequential

Episodic:

- In episodic environments, the agent's experience is **divided into atomic episodes**.
 - Each episode consists of the agent perceiving and then performing a single action.
 - The next episode **does not depend** on the actions taken in previous episodes.
- Example is the classification tasks

Sequential:

- In sequential environments, **the current decision could affect all future decisions**.
- Examples are Chess and taxi driving

Environment Types

Static vs. Dynamic

Static:

- **Static environments are unchanged** and easy to deal with because the agent need not keep looking at the world while it is deciding on an action.
- Crossword puzzles are static.

Dynamic:

- If the **environment can be changed** while an agent is deliberating, then we say the environment is dynamic for that agent --- taxi driving
- ▶ *If the environment itself does not change with the passage of time but the agent's performance score changes*, then we say the environment is **semi-dynamic** --- Chess when played with a clock, is semi-dynamic

Environment Types

Discrete vs. Continuous

- ▶ The discrete/continuous distinction can be applied to the **state of the environment**, to the way **time** is handled, and to the **percept** and **actions**.
 - **Chess** has a **discrete set of percept and actions**.
 - **Taxi driving** contains a **continuous state** and **continuous-time problem**,
 - **Taxi-driving** actions are also **continuous**.

Environment Types

Single agent vs. multi-agent

- ▶ An agent operating by itself in an environment is a single agent.
- ▶ Examples:
 - Crossword puzzle -> a single agent
 - chess -> two-agents.

Does an agent “A” have to treat an object “B” as an agent, or can it be treated merely as a stochastically behaving object

- The key distinction is whether B's behaviour is best described as **maximizing a performance measure** whose value **depends on agent A's behaviour**.

Agent functions and Programs

- ▶ **Agent program**

takes the **current percept** as an input from the sensors and returns an action to the actuators.

- ▶ **Agent function**

takes the **whole percept history** and maps onto actions.

- ▶ **Notice the difference** between the agent program, which takes the current percept as input, and the agent function, which takes the entire percept history.
- ▶ The agent needs to remember the whole percept sequence, if it needs it.

Table-driven Agent

Table-driven Agent Program

- ▶ A trivial agent program **keeps track of the percept sequence** and then uses it to index into a table of actions to decide what to do

function TABLE-DRIVEN-AGENT(*percept*) **returns** an action

persistent: *percepts*, a sequence, initially empty

table, a table of actions, indexed by percept sequences, initially fully specified

append *percept* to the end of *percepts*

action \leftarrow LOOKUP(*percepts*, *table*)

return *action*

Percept sequence	Actions
[A,Clean]	Right
[A, Dirty]	Suck
[B,Clean]	Left
[B,Dirty]	Suck
[A,Clean],[A,Clean]	Right
[A,Clean],[A,Dirty]	Suck
...	...
[A,Clean],[A,Clean],[A,Clean]	Right
[A,Clean],[A,Clean],[A, Dirty]	Suck

Table-driven Agent

Why the table-driven approach for agent construction is considered as failure.

- ▶ The lookup table will contain the **number of entries**

$$\sum_{t=1}^T |\mathcal{P}|^t$$

Where,

- ▶ \mathcal{P} is the set of percept
- ▶ T is the lifetime

Table-driven Agent

Example 1: Automated taxi:

- ▶ The visual input from a single camera comes in at the rate of roughly **27 megabytes per second** with info
 - ❑ 30 frames per second,
 - ❑ 640 x 480 pixels with 24 bits of colour information
- ▶ This gives a lookup table with over $10^{250,000,000,000}$ entries for an hour's driving.

Example 2: Chess:

- ▶ Even the lookup table for chess—a tiny, well-behaved real world—would have at least 10^{150} entries.

Agent Types



Agent Types

- ▶ There are following four kinds of agent
 - ❑ Simple reflex agents
 - ❑ Model-based reflex agents
 - ❑ Goal-based agents
 - ❑ Utility-based agents

Simple Reflex Agents

- ▶ Select actions on the basis of the **current percept** and ignoring the rest of the percept history


- ▶ **Condition-action rule**

if car-in-front-is-braking then initiate-braking.

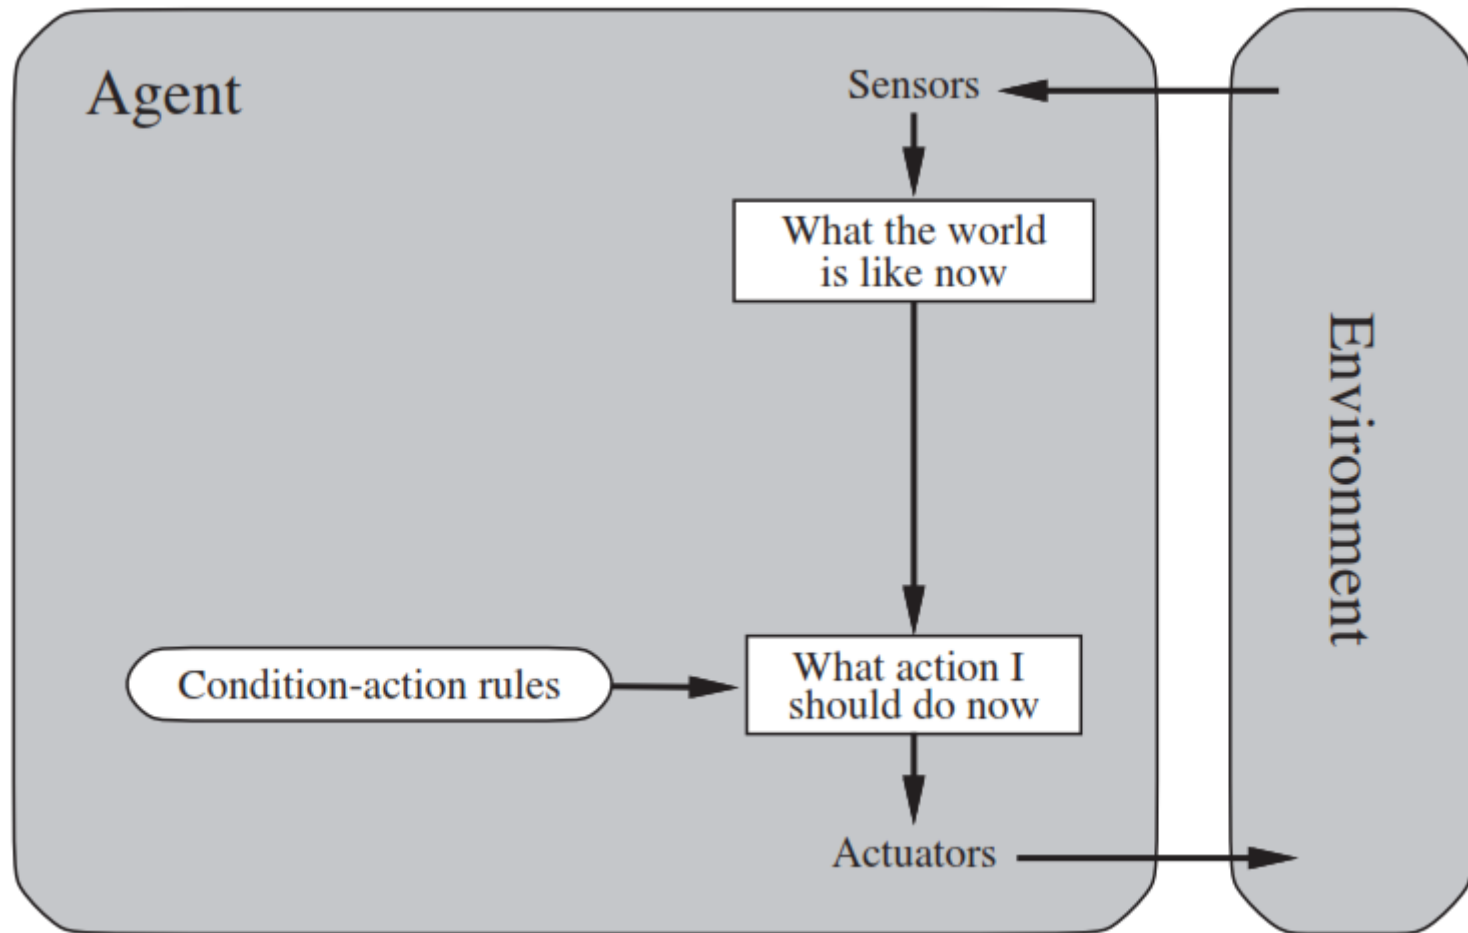
- ▶ **Vaccum Cleaner World:**

function REFLEX-VACUUM-AGENT(*[location,status]*) **returns** an action

if status = Dirty then return Suck
else if location = A then return Right
else if location = B then return Left



Simple Reflex Agents



Simple Reflex Agents

```
function SIMPLE-REFLEX-AGENT(percept) returns an action  
  persistent: rules, a set of condition–action rules  
  
  state  $\leftarrow$  INTERPRET-INPUT(percept)  
  rule  $\leftarrow$  RULE-MATCH(state, rules)  
  action  $\leftarrow$  rule.ACTION  
  return action
```

- ▶ The agent will work only if
 - the **correct decision** can be made on the basis of the **current percept** –that is, only if the environment is **fully observable**.

Model-based Reflex Agents

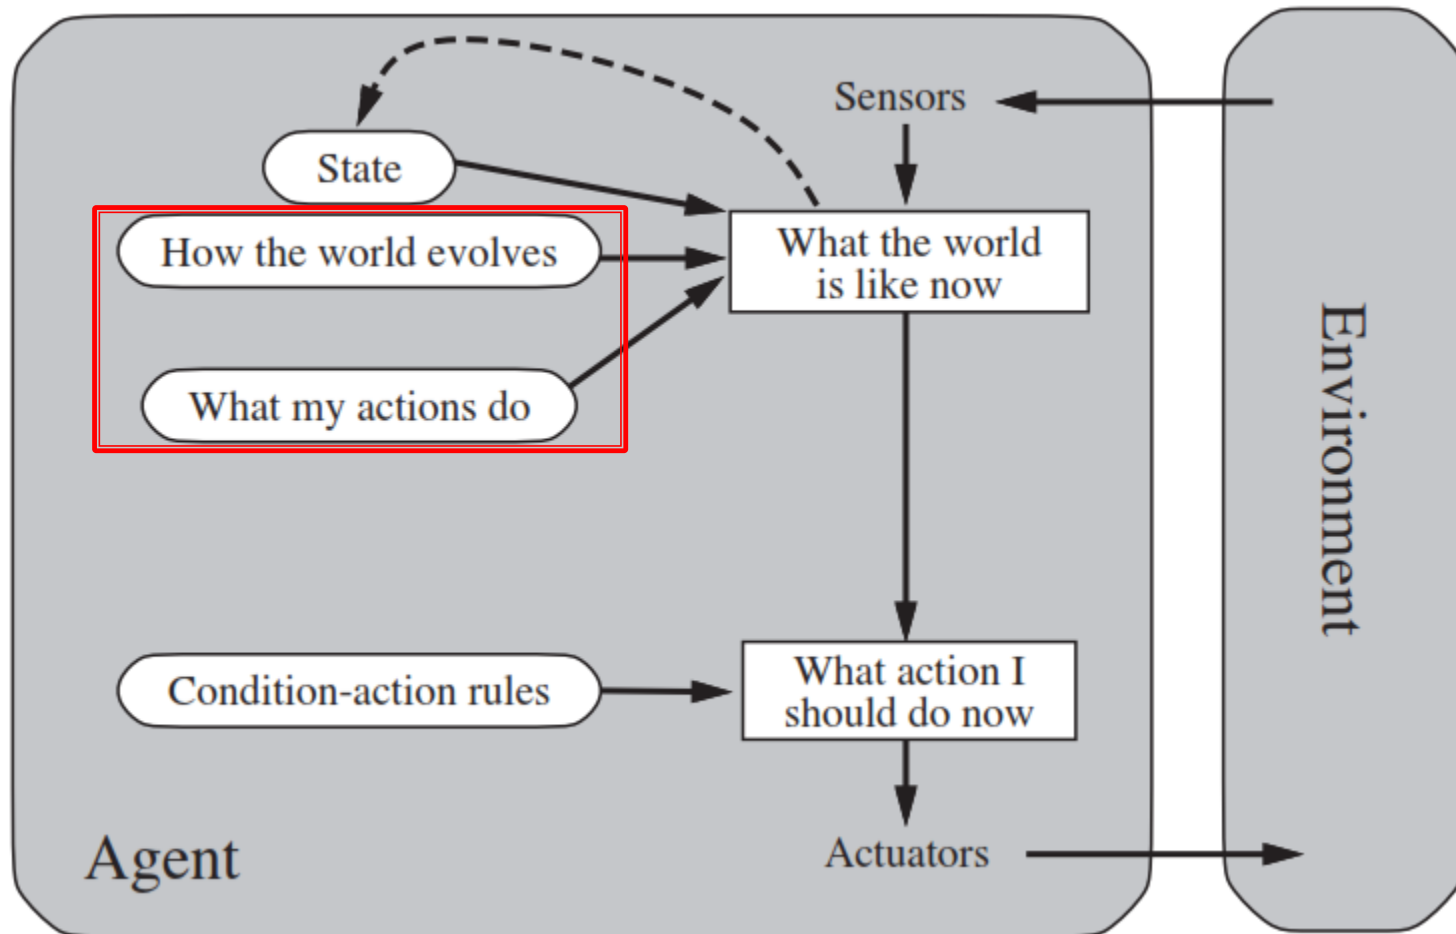
Model:

- ▶ *A description that how the next state depends on the current state & action.*
- ▶ It **handles partial observability** in a more effective way.
- ▶ It maintains some sort of **internal state** that depends on the percept history and thereby **reflects at least some of the unobserved aspects of the current state.**

Model-based Reflex Agents

- ▶ Updating this internal state information requires **two kinds of knowledge**:
 - ❑ **First**, how the **world evolves** independently of the agent.
 - ❑ **Second**, how the **agent's own actions** affect the world.

Model-based Reflex Agents



Model-based Reflex Agents

```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
  persistent: state, the agent's current conception of the world state
               model, a description of how the next state depends on current state and action
               rules, a set of condition–action rules
               action, the most recent action, initially none

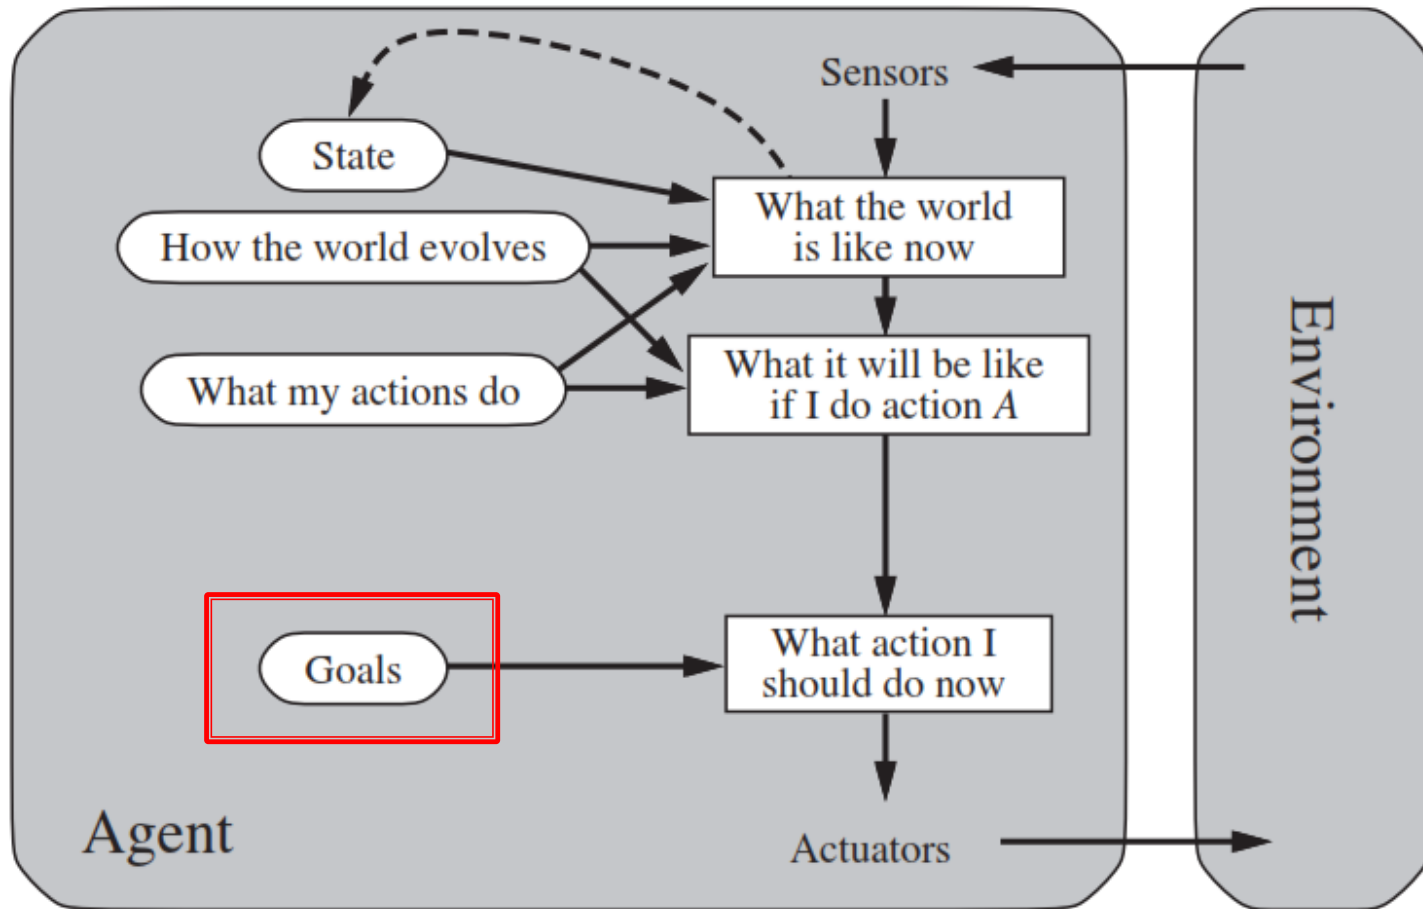
  state ← UPDATE-STATE(state, action, percept, model)
  rule ← RULE-MATCH(state, rules)
  action ← rule.ACTION
  return action
```

The **Model** is the knowledge about “how the world works”.

Goal-based Agents

- ▶ Information about the **current state of the environment is not always enough** to decide what to do (e.g. **decision at a road junction**).
- ▶ The agent needs some sort of **goal information** that describes situations that are desirable.
- ▶ The agent program can combine this with information about the **results of possible actions** in order to choose actions that achieve the goal.
- ▶ Usually requires some **search and planning**.

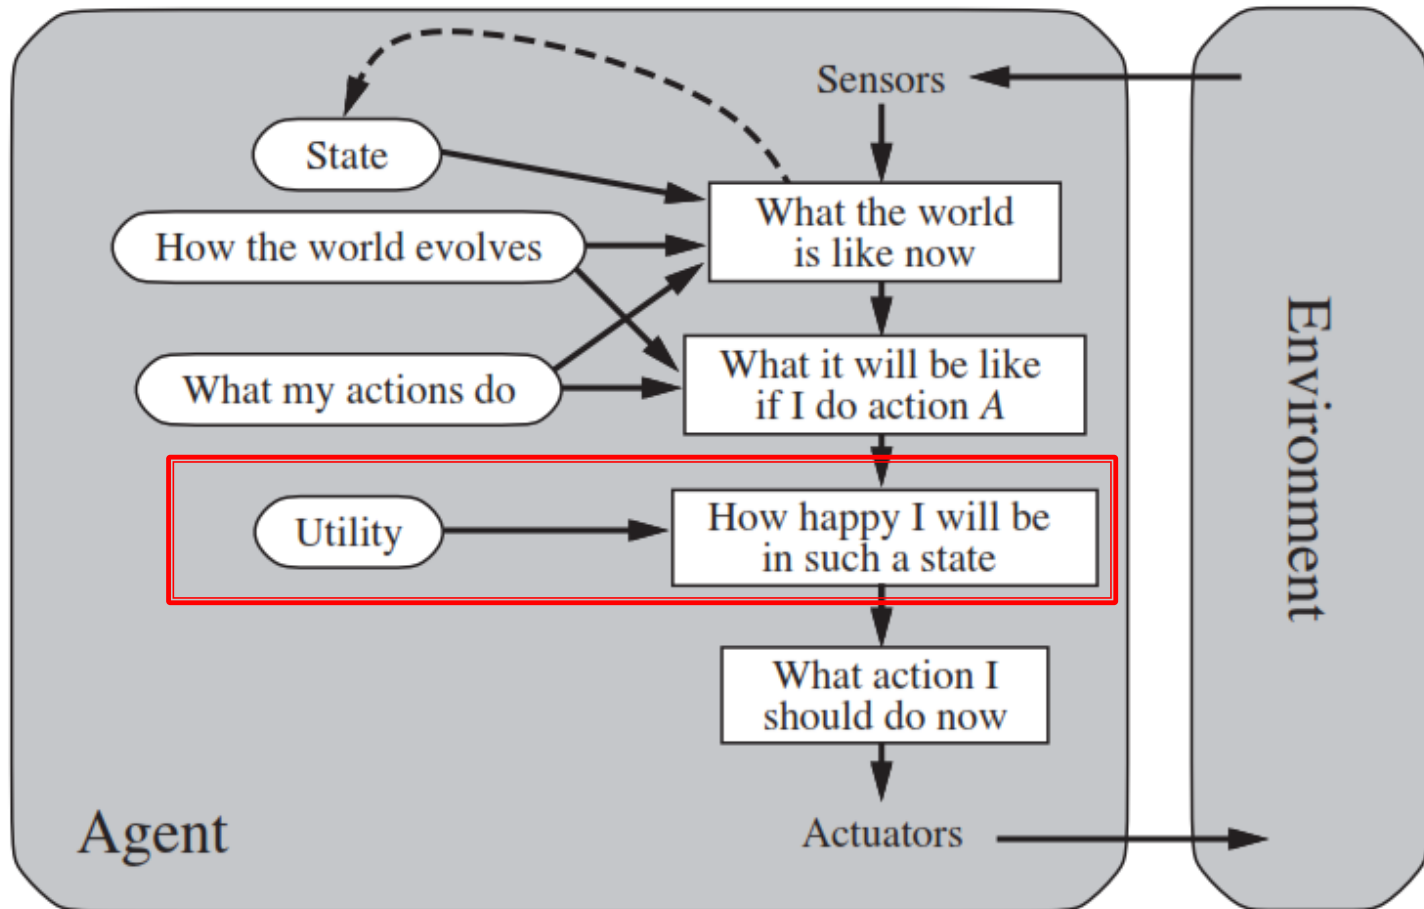
Goal-based Agents



Utility-based Agents

- ▶ Goals provide a **binary distinction between happy and unhappy states.**
- ▶ Agents so far we have discussed had a single goal. Agents may have to *juggle conflicting goals.*
- ▶ Need to optimise utility over a range of goals.
- ▶ **Utility**: measure of *happiness (a real number)*, --- *the quality of being useful .*
- ▶ A **utility function** *maps a state onto a real number which describes the associated degree of happiness.*

Utility-based Agents



Reading Material

- ▶ **Artificial Intelligence, A Modern Approach**
Stuart J. Russell and Peter Norvig
 - Chapter 2.

