

AI 2002

Artificial Intelligence

Dr. Hashim Yasin

Adversarial Search

Adversarial Search

- ▶ **Competitive** environments, in which the **agents' goals are in conflict**, giving rise to **adversarial search** problems—often known as **games**

Why do AI researchers study game playing?

- ▶ It's a **good reasoning problem**, formal and nontrivial.
- ▶ Offer an opportunity to study problems involving **{hostile, adversarial, competing} agents**.
- ▶ **Direct comparison** with humans and other computer programs is easy.

Adversarial Search

Mainly games of strategy with the following characteristics:

- ▶ Sequence of **moves** to play
- ▶ Rules that specify **possible moves**
- ▶ Rules that specify a **payment** for each move
- ▶ Objective is to **maximize** your payment

Games

▶ **Competitive: Commonly Zero Sum**

- One player wins and the other loses
- A zero-sum game is defined as one where the **total payoff to all players is the same for every instance** of the game. Chess is zero-sum because every game has payoff of either $0 + 1$, $1 + 0$ or $\frac{1}{2}, \frac{1}{2}$.

▶ **Perfect Information:**

- Players know the results of all previous moves
- There is one best way to win the game for all players

▶ **Imperfect Information**

- Players do not know all of the previous moves

Games

- ▶ **Initial State (s_0):** The initial state, which specifies how the game is *set up at the start*.
- ▶ **Players:** defines which player has the move in a state.
- ▶ **Actions:** The set of legal moves.
- ▶ **Result (s, a):** The transition model, which defines the result of a move. The state after action **a** is the state **s** .
- ▶ **Terminal Test:** A terminal test, which is true when the game is over and false otherwise.


Games

- ▶ **Terminal State:** States where the game has ended are called terminal states.
- ▶ **Utility:** A utility function (also called an **objective function** or **payoff function**), defines the final numeric value for a game that ends in terminal state ***s*** for a player ***p***.
 - In chess, the outcome is a win, loss, or draw, with values +1, 0, or $\frac{1}{2}$.
 - Some games have a wider variety of possible outcomes; the payoffs in **backgammon** range from 0 to +192.

Game Tree

- ▶ The **initial state**, **actions**, and **results** define the game tree for the game.
- ▶ A **game tree** where the **nodes are the game states** and **the edges are moves**.
- ▶ The game tree is best thought of as a theoretical construct that we cannot realize in the physical world.
 - For **tic-tac-toe** the game tree is relatively small—fewer than $9! = 362,880$ terminal nodes.
 - For **chess** there are over 10^{40} nodes,

Game Tree: tic-tac-toe



Minimax

- ▶ Minimax is a method used **to evaluate game trees**.
- ▶ Given a game tree, the optimal strategy can be determined from the **minimax value** of each node.
- ▶ *A static evaluator is applied to leaf nodes, and **the values are passed back up the tree to determine the best score** the computer can obtain against a rational opponent.*

Minimax


MAX

- ▶ Wants to **maximize** the result of the utility function
- ▶ **Winning strategy** if, on MIN's turn, a win is obtainable for MAX for all moves that MIN can make

MIN

- ▶ Wants to **minimize** the result of the utility function
- ▶ **Winning strategy** if, on MAX's turn, a win is obtainable for MIN for all moves that MAX can make


Minimax



MINIMAX(s) =

$$\begin{cases}
 \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\
 \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\
 \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN}
 \end{cases}$$


Minimax



MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax



MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax

function MINIMAX-DECISION(*state*) *returns an action*
 return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(s, a))$

function MAX-VALUE(*state*) *returns a utility value*
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$
 return *v*

function MIN-VALUE(*state*) *returns a utility value*
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow \infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$
 return *v*

Properties of Minimax

Complete?

- ▶ Yes (if tree is finite).

Optimal?

- ▶ Yes

Time complexity?

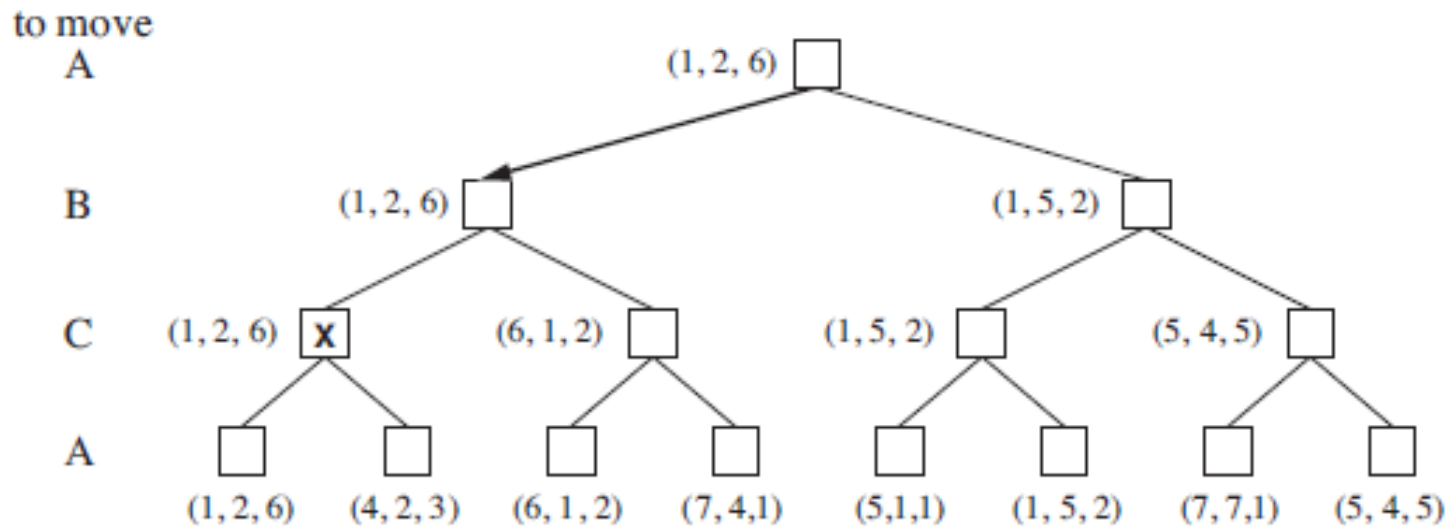
- ▶ $O(b^m)$, m is the maximum depth of the tree and b is the legal moves.

Space complexity?

- ▶ $O(bm)$
 - (depth-first search, generate all actions at once)
- ▶ $O(m)$
 - (backtracking search, generate actions one at a time)

Multiplayer Games

- ▶ Each node must hold **a vector of values**
- For example**, for three players A, B, C (V_A, V_B, V_C)
- ▶ The backed up vector at node n will *always be the one that maximizes the payoff* of the player choosing at n




Searching Game Trees

- ▶ **Exhaustively searching** a game tree is not usually a good idea.
- ▶ Even for a simple game like
 - **tic-tac-toe** there are over **350,000 nodes** in the complete game tree.
- ▶ **An additional problem** is that the computer only gets to choose every other path through the tree and the opponent chooses the others.

Pruning

MAX

MIN




$$\begin{aligned} \text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \end{aligned}$$

Pruning

MAX

MIN



$$\begin{aligned}
 \text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\
 &= \max(3, \min(2, x, y), 2) \\
 &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\
 &= 3.
 \end{aligned}$$

Do we need z ?

Pruning

- ▶ We can use a **branch-and-bound technique** to reduce the number of states that must be examined to determine the value of a tree.

Branch-and-bound Technique:

- ▶ We keep track of a **lower bound on the value of a maximizing node**, and don't bother evaluating any trees that cannot improve this bound.
- ▶ Keep track of an **upper bound on the value of a minimizing node**. Don't bother with any sub-trees that cannot improve this bound.

Minimax with Alpha-Beta Cutoffs

Alpha Cutoffs:

- ▶ Alpha is the *lower bound on maximizing nodes*.


Beta Cutoffs:

- ▶ Beta is the *upper bound on minimizing nodes*.
- ▶ Both alpha and beta get passed down the tree during the Minimax search.


Minimax with Alpha-Beta Cutoffs

- ▶ At minimizing nodes, we stop evaluating children if we get a child whose value is **less than the current lower bound (*alpha*)**.
- ▶ At maximizing nodes, we stop evaluating children as soon as we get a child whose value is **greater than the current upper bound (*beta*)**.
- ▶ Some branches will never be played by rational players since they include sub-optimal decisions (for either player)


Alpha-Beta Pruning: Example




Alpha-Beta Pruning: Example



Alpha-Beta Pruning: Example




Alpha-Beta Pruning: Example



Alpha-Beta Pruning: Example

At MIN we will change
the value of BETA


At MAX we will change
the value of ALPHA



Alpha-Beta Pruning: Example

At MIN we will change
the value of BETA


At MAX we will change
the value of ALPHA



Alpha-Beta Pruning: Example

At MIN we will change
the value of BETA

At MAX we will change
the value of ALPHA



Alpha-Beta Pruning: Effectiveness

- ▶ The effectiveness **depends on the order** in which children are visited.
- ▶ **In the best case**, the effective branching factor will be reduced from **b** to **\sqrt{b}** .
- ▶ **In an average case** (random values of leaves) the branching factor is reduced to $\frac{b}{\log b}$.

Reading Material

- ▶ **Artificial Intelligence, A Modern Approach**
Stuart J. Russell and Peter Norvig
 - Chapter 5.

