

AI 2002

Artificial Intelligence

Dr. Hashim Yasin

Hill-climbing Search

Hill-climbing Search

- ▶ It is simply a **loop** that continuously moves in the direction of increasing value—that is, uphill.
- ▶ It terminates when it reaches a “peak” where no neighbor has a higher value.
 - *does not maintain a search tree*
 - need only
 - *record the state* and
 - the value of the *objective function*.
- ▶ Hill climbing only looks towards the **immediate neighbors** of the current state.

Hill-climbing Search

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

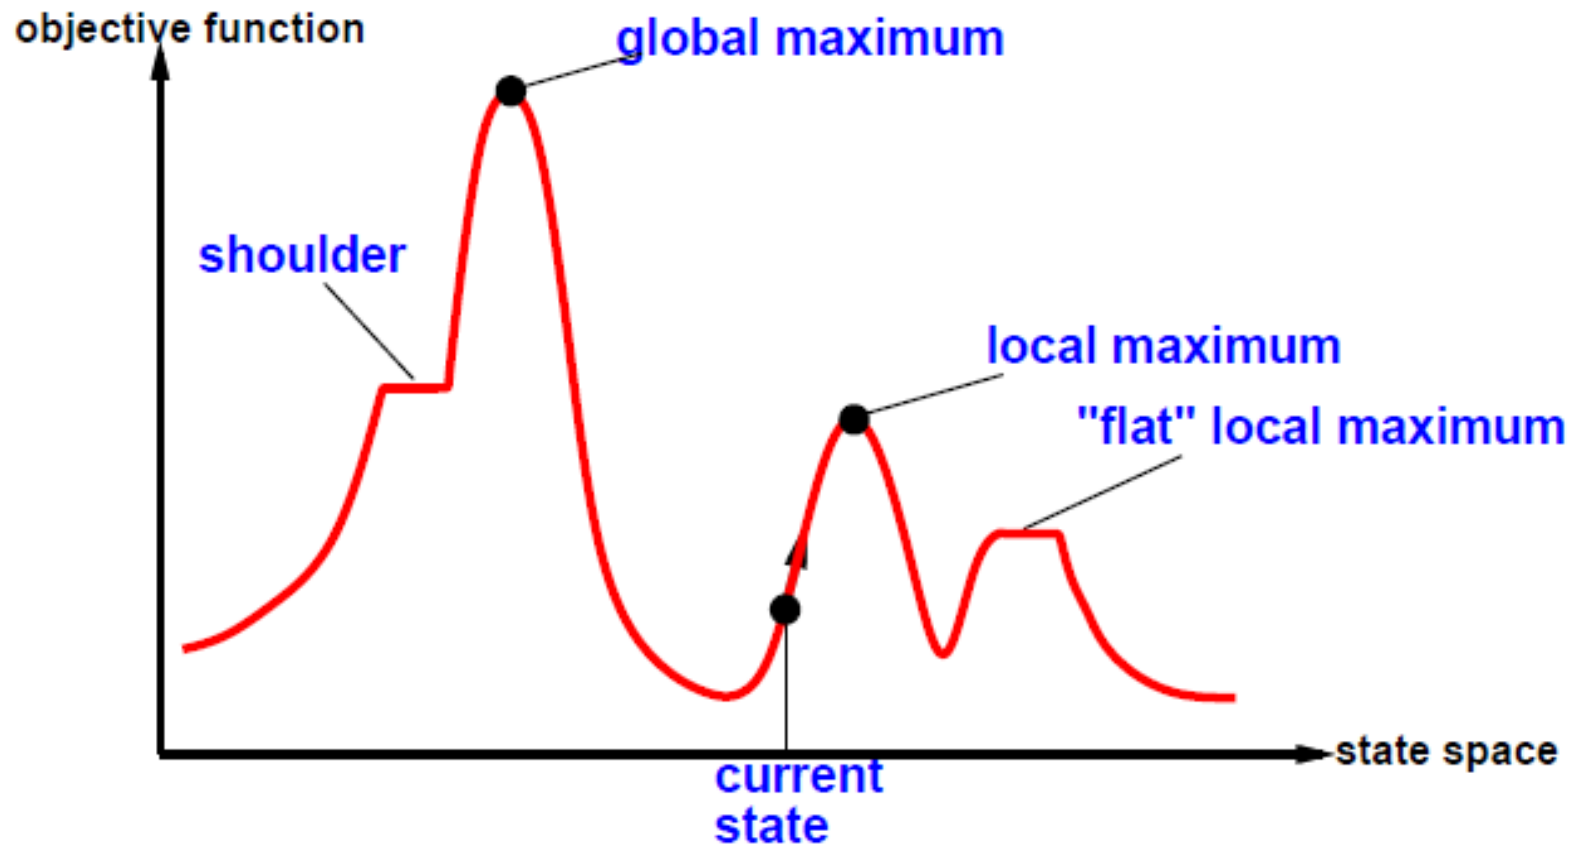
neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

- ▶ At *each step the current node is replaced by the best neighbor*; the neighbor with the highest VALUE,
- ▶ If a **heuristic cost estimate h** is used, we would find the neighbor with the ***lowest h*** .

Hill-climbing Search



Hill-Climbing Search

With randomly generated 8-queens starting states, the steepest-ascent hill climbing:

- ▶ 14% of the time it solves the problem
- ▶ 86% of the time it get stuck at a local minimum
- ▶ However...
 - Takes only 4 steps on average when it succeeds
 - And 3 on average when it gets stuck
 - (for a state space with $8^8 \approx 17$ million states)

Hill-Climbing: Variants

Stochastic hill-climbing:

- ▶ Chooses randomly among *potential* successors
- ▶ Sometimes better than steepest ascent

First-choice hill-climbing:

- ▶ Generates successors randomly and picks first
- ▶ Good for many successors

Random restart hill-climbing:

- ▶ Restarts from randomly generated initial state when failed
- ▶ Roughly 7 iterations with 8-queens problem

Hill-Climbing Search

- ▶ **The success of hill-climbing search** depends very much
 - on the *shape of the state-space landscape*
- ▶ If there are few local maxima and plateaux,
 - **random-restart hill climbing** will find a good solution very quickly.

Simulated Annealing

Simulated Annealing

Idea:

- ▶ escape local maxima by allowing some “bad” moves but gradually *decrease the size and frequency of the bad moves*,
- ▶ In thermodynamics, the probability to go from a state with energy E_1 to a state of energy E_2 is given by:

$$p = e^{\frac{(E_2 - E_1)}{kT}} = e^{\frac{-(E_1 - E_2)}{kT}}$$

Simulated Annealing

$$p = e^{\frac{-(E_1 - E_2)}{kT}}$$

Where,

- ▶ e is Euler's number
- ▶ T is a "temperature" controlling the probability of downward steps
- ▶ k is Boltzmann's constant
 - (relating energy and temperature; with appropriate choice of units, it will be equal to 1).

Simulated Annealing

$$p = e^{\frac{(E_2 - E_1)}{kT}} = e^{\frac{-(E_1 - E_2)}{kT}}$$

- ▶ The idea is that probability decreases exponentially with $E_2 - E_1$ increasing,
- ▶ The probability gets lower as temperature decreases
- ▶ If the *schedule* lowers T slowly enough, the algorithm will find a global optimum with probability approaching 1.

Simulated Annealing

function **SIMULATED-ANNEALING**(*problem*, *schedule*) returns a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

local variables: *current*, a node

next, a node

T, a “temperature” controlling prob. of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 to ∞ do

T \leftarrow *schedule*[*t*]

if *T* = 0 then return *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] – VALUE[*current*]

if $\Delta E > 0$ then *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

Similar to hill climbing,
but a random move instead
of best move

case of improvement, make
the move

Otherwise, choose the move with probability that
decreases exponentially with the “badness” of the move.

Simulated Annealing...Example

- ▶ Consider there are 3 moves available, with changes in the objective function of

$$\Delta E_1 = -0.1, \Delta E_2 = 0.5, \Delta E_3 = -3$$

- ▶ Suppose $T = 1$
- ▶ Pick a move randomly:
 - if ΔE_2 is picked, move there.
 - if ΔE_1 or ΔE_3 are picked, probability of move = $e^{\frac{\Delta E}{T}}$
 - move 1: prob1 = $e^{-0.1} = 0.9$,
 - i.e., 90% of the time we will accept this move
 - move 3: prob3 = $e^{-3} = 0.0497$
 - i.e., 5% of the time we will accept this move

Simulated Annealing

T = “temperature” parameter

- ▶ **If T is high** \Rightarrow the probability of “locally bad” move is higher
- ▶ **If T is low** \Rightarrow the probability of “locally bad” move is lower
- ▶ typically, T is decreased as the algorithm runs longer
 - i.e., there is a “temperature schedule”

Simulated Annealing

Convergence:

- ▶ With exponential schedule, will provably converge to global optimum
 - If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1.
- ▶ Few more precise convergence rate
 - ▶ Recent work on **rapidly mixing Markov chains**.
 - ▶ Surprisingly, deep foundations.

Simulated Annealing

- ▶ method proposed in 1983 by IBM researchers for solving VLSI layout problems.
 - theoretically will always find the global optimum (the best solution)
- ▶ Useful for some problems, but can be very slow
 - slowness comes about because T must be decreased very gradually to retain optimality
- ▶ In practice how to decide the rate at which to decrease T ? (this is a practical problem with this method)

Local Beam Search

Local Beam Search

Idea:

- ▶ Keep track of k states rather than just one
- ▶ Start with k randomly generated states
- ▶ At each iteration, all the successors of all k states are generated
- ▶ If any one is a goal state, stop;
- ▶ else select the k best successors from the complete list and repeat.

Local Beam Search

```
function BEAM-SEARCH(problem, k) returns a solution state
  start with k randomly generated states
  loop
    generate all successors of all k states
    if any of them is a solution then return it
    else select the k best successors
```

A local beam search with k states might seem to be nothing more than running **k random restarts in parallel** instead of in sequence.

Local Beam Search

Local beam search with $k = 1$

- ▶ We would randomly generate **1 start state**
- ▶ At each step we would generate all the successors, and retain the **1 best state**
- ▶ Equivalent to **HILL-CLIMBING**

Local Beam Search

Local beam search with $k = \infty$

- ▶ **1 initial state** and no limit of the number of states retained
- ▶ We start at initial state and **generate ALL successor** states (no limit how many)
- ▶ If one of those is a goal, we stop
- ▶ Otherwise, we generate all successors of those states (2 steps from the initial state), and continue
- ▶ Equivalent to **BREADTH-FIRST SEARCH** except that *each layer is generated all at once.*

Reading Material

- ▶ **Artificial Intelligence, A Modern Approach**
Stuart J. Russell and Peter Norvig
 - Chapter 4.

