

AI 2002

Artificial Intelligence

Dr. Hashim Yasin

Perceptron Training Rule

- ▶ The ***perceptron training rule***, which revises the weight w_i associated with input x_i according to the rule:

$$w_i \leftarrow w_i + \Delta w_i$$

where

$$\Delta w_i = \eta(t - o)x_i$$

Where:

- t is target value
- o is perceptron output
- η is small constant (e.g., 0.1) called *learning rate*



Perceptron Training Rule

- ▶ The **perceptron rule** finds a successful weight vector when the training examples are **linearly separable**,
- ▶ It fails to converge if the examples are **not linearly separable**.
- ▶ The solution is ... **Delta Rule** also known as (**Widrow-Hoff Rule**)

Delta Rule

- ▶ use ***gradient descent*** to search the hypothesis space of possible weight vectors to find the weights that best fit the training examples.

Perceptron



The **decision surface** represented by a **two-input perceptron x_1 and x_2** .
(a) A set of training examples and the decision surface of a perceptron that classifies them correctly. (b) A set of training examples that is not linearly separable.

Delta Rule

Delta Rule

- ▶ In perceptron training rule we employ *thresholded perceptron*

$$o(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} > 0 \\ -1 & \text{otherwise} \end{cases}$$

- ▶ The delta training rule is the task of training an *unthresholded perceptron*;
 - a *linear unit* for which the output o is given by

$$o(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$$

- A linear unit corresponds to the *first stage* of a perceptron, *without* the threshold.

Delta Rule

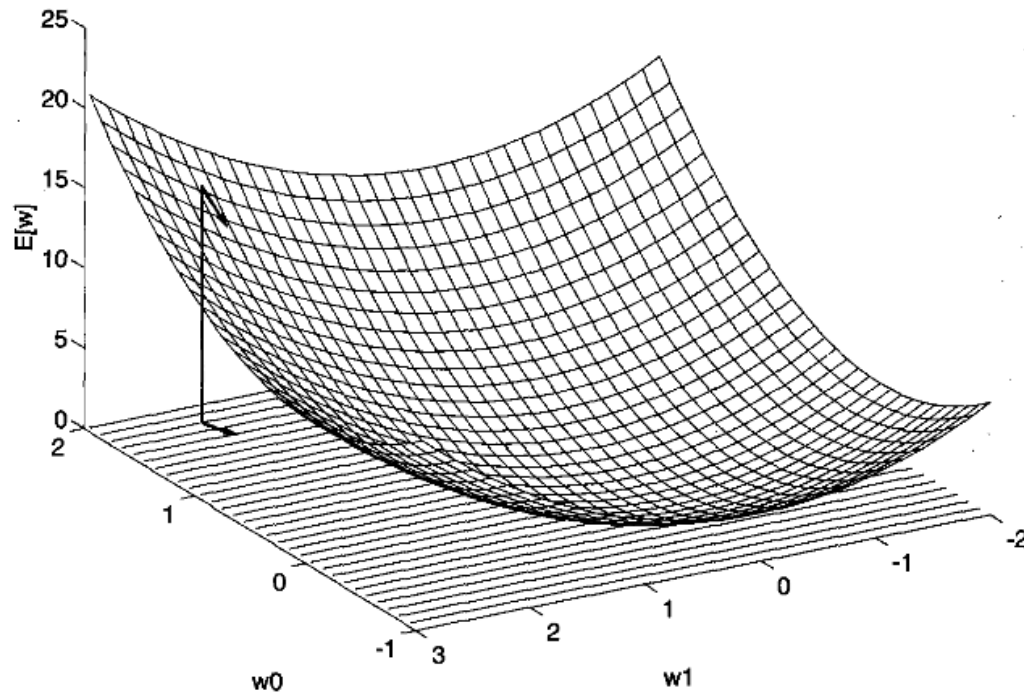
- ▶ In order to derive **a weight learning rule for linear units**,
 - Specify a measure for the **training error** of a hypothesis (weight vector), relative to the training examples.

$$E(\vec{w}) \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

- D is the set of training examples,
- t_d is the target output for training example d ,
- o_d is the output of the linear unit for training example d
- E is characterized as a function of \mathbf{w} , because the linear unit output o depends on this weight vector.

Hypothesis Space

- ▶ For a linear unit with two weights, the hypothesis space H is the w_0, w_1 plane.



Error of different hypotheses.

Gradient Descent

- ▶ Gradient descent search determines *a weight vector that minimizes E* by
 - ❑ Starting with an arbitrary initial weight vector,
 - ❑ Repeatedly modifying it in small steps.
 - ❑ At each step, the ***weight vector is altered in the direction*** that produces the **steepest descent** along the error surface,
 - ❑ This process continues until the **global minimum error** is reached.

Gradient Descent

How can we calculate the direction of steepest descent along the error surface?

- ▶ This direction can be found by computing the derivative of E with respect to each component of the vector \mathbf{w} .

$$\nabla E(\vec{w}) \equiv \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

- ▶ Notice $\nabla E(\vec{w})$ is itself a vector, whose components are the partial derivatives of E with respect to each of the w_i .

Gradient Descent

- ▶ The **gradient specifies the direction** that produces the steepest increase in E . The training rule for gradient descent is,

$$\vec{w} \leftarrow \vec{w} + \Delta \vec{w}$$

where,

$$\Delta \vec{w} = -\eta \nabla E(\vec{w})$$

- ▶ The negative sign is present because we want to move the weight vector in the direction that ***decreases*** E .

Gradient Descent

- ▶ This training rule can also be written in its **component form**,

$$\vec{w} \leftarrow \vec{w} + \Delta \vec{w}$$

$$w_i \leftarrow w_i + \Delta w_i$$

where,

$$\Delta \vec{w} = -\eta \nabla E(\vec{w})$$

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

- ▶ The steepest descent is achieved by altering each component w_i in proportion to $\frac{\partial E}{\partial w_i}$

Gradient Descent

- ▶ The vector of derivatives $\frac{\partial E}{\partial w_i}$ that form the gradient can be obtained by differentiating E from delta rule

$$E(\vec{w}) \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

$$o(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$$

$$\begin{aligned} \frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_{d \in D} \frac{\partial}{\partial w_i} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_{d \in D} 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d) \\ &= \sum_{d \in D} (t_d - o_d) \frac{\partial}{\partial w_i} (t_d - \vec{w} \cdot \vec{x}_d) \end{aligned}$$

$$\frac{\partial E}{\partial w_i} = \sum_{d \in D} (t_d - o_d) (-x_{id})$$

Gradient Descent

- ▶ We now have an equation that gives $\frac{\partial E}{\partial w_i}$ in terms of
 - the linear unit inputs x_{id} ,
 - outputs o_d ,
 - target values t_d associated with the training examples
- ▶ The weight update rule for gradient descent becomes,

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

$$\Delta w_i = \eta \sum_{d \in D} (t_d - o_d) x_{id}$$

Gradient Descent

GRADIENT-DESCENT(*training_examples*, η)

Each training example is a pair of the form $\langle \vec{x}, t \rangle$, where \vec{x} is the vector of input values, and t is the target output value. η is the learning rate (e.g., .05).

- Initialize each w_i to some small random value
- Until the termination condition is met, Do
 - Initialize each Δw_i to zero.
 - For each $\langle \vec{x}, t \rangle$ in *training_examples*, Do
 - Input the instance \vec{x} to the unit and compute the output o
 - For each linear unit weight w_i , Do

$$\Delta w_i \leftarrow \Delta w_i + \eta(t - o)x_i$$

- For each linear unit weight w_i , Do

$$w_i \leftarrow w_i + \Delta w_i$$

$$\Delta w_i = \eta \sum_{d \in D} (t_d - o_d) x_{id}$$

Gradient Descent

- ▶ Gradient descent is an important general paradigm for learning.
- ▶ It is a *strategy for searching through a large or infinite hypothesis space* that can be applied whenever
 - 1) the hypothesis space contains **continuously parameterized hypotheses** (e.g., the weights in a linear unit),
 - 2) the **error can be differentiated** with respect to these hypothesis parameters

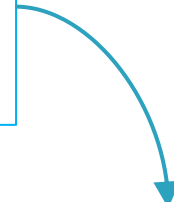
Gradient Descent

- ▶ The **key practical difficulties** in applying gradient descent are:
 - a) **converging to a local minimum*** can sometimes be **quite slow** (i.e., it can require many thousands of gradient descent steps),
 - b) **if there are multiple local minima*** in the error surface, then there is no guarantee that the procedure will find the global minimum.

Stochastic Gradient Descent

- ▶ The idea behind stochastic gradient descent is to approximate the gradient descent search by **updating weights incrementally**, following the calculation of the error for **each individual example**.

$$\Delta w_i = \eta \sum_{d \in D} (t_d - o_d) x_{id}$$


$$\Delta w_i = \eta (t - o) x_i$$

Stochastic Gradient Descent

- ▶ One way to view this stochastic gradient descent is to consider a **distinct error function** defined for *each individual training example d as follows*.

$$E_d(\vec{w}) = \frac{1}{2}(t_d - o_d)^2$$

- where t , and o_d are the target value and the unit output value for training example d .
- Stochastic gradient descent iterates over the training examples d in D ,
- at each iteration altering weights according to the gradient

Stochastic Gradient Descent

GRADIENT-DESCENT(*training_examples*, η)

Each training example is a pair of the form $\langle \vec{x}, t \rangle$, where \vec{x} is the vector of input values, and t is the target output value. η is the learning rate (e.g., .05).

- Initialize each w_i to some small random value
- Until the termination condition is met, Do
 - For each $\langle \vec{x}, t \rangle$ in *training_examples*, Do
 - Input the instance \vec{x} to the unit and compute the output o
 - For each linear unit weight w_i , Do

$$w_i \leftarrow w_i + \eta(t - o)x_i$$

The Key Difference

- ▶ In **stochastic gradient descent**, weights are updated upon examining *each* training example.
- ▶ Whereas in **standard gradient descent**, the error is summed over *all* examples before updating weights,
 - ❑ **Standard gradient descent** requires **more computation** per weight update step.
 - ❑ **Standard gradient descent** is often used with a **larger step size** per weight update than stochastic gradient descent.

The Key Difference

- ▶ When there are multiple local minima with respect to $E(\mathbf{w})$,
 - ❑ The **stochastic gradient descent** can sometimes *avoid falling into these local minima*,
 - ❑ It is due to the reason that it uses various $\nabla E_d(\vec{\mathbf{w}})$ rather than $\nabla E(\vec{\mathbf{w}})$ to guide its search.


Training Rules

- ▶ **Perceptron Training Rule:** guarantee to succeed if
 - training examples are linearly separable
 - Sufficiently small learning rate
- ▶ **Delta Rule:**
 - use gradient descent
 - converges only asymptotically toward the minimum error hypothesis,
 - converges regardless of whether the training data are linearly separable.

Multilayer Networks

Multilayer Perceptron Architecture

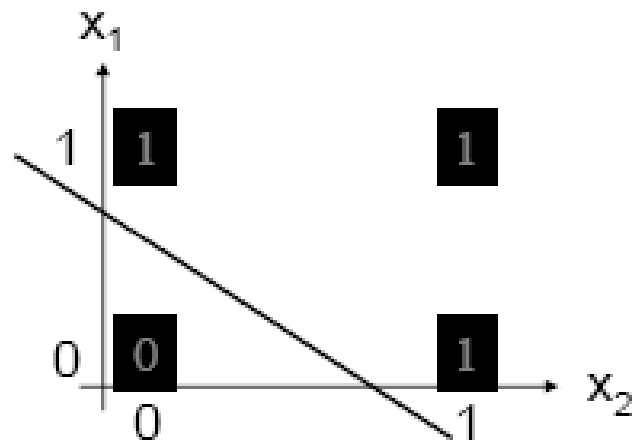
MLP used to describe any general feedforward (no recurrent connections) network



Multilayer Networks

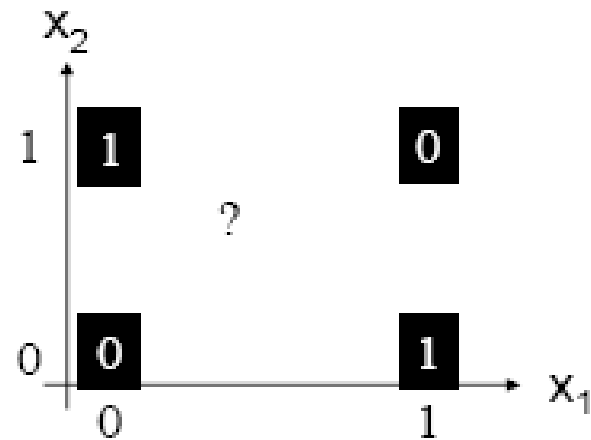
OR function

x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	1




XOR function

x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0



Multilayer Networks



Network Topology:

2 hidden nodes

1 output

Weights:

$$w_{11} = w_{12} = 1$$

$$w_{21} = w_{22} = 1$$

$$w_{01} = -1.5$$

$$w_{02} = -0.5$$

$$w_{13} = -1$$

$$w_{23} = 1$$


$$w_{03} = -0.5$$

Desired behavior:

x_1	x_2	o_1	o_2	y
0	0			
1	0			
0	1			
1	1			

Piecewise linear classification using an MLP with threshold (perceptron) units

Multilayer Networks



Network Topology:

2 hidden nodes

1 output

Weights:

$$w_{11} = w_{12} = 1$$

$$w_{21} = w_{22} = 1$$

$$w_{01} = -1.5$$

$$w_{02} = -0.5$$

$$w_{13} = -1$$

$$w_{23} = 1$$

$$w_{03} = -0.5$$

Desired behavior:

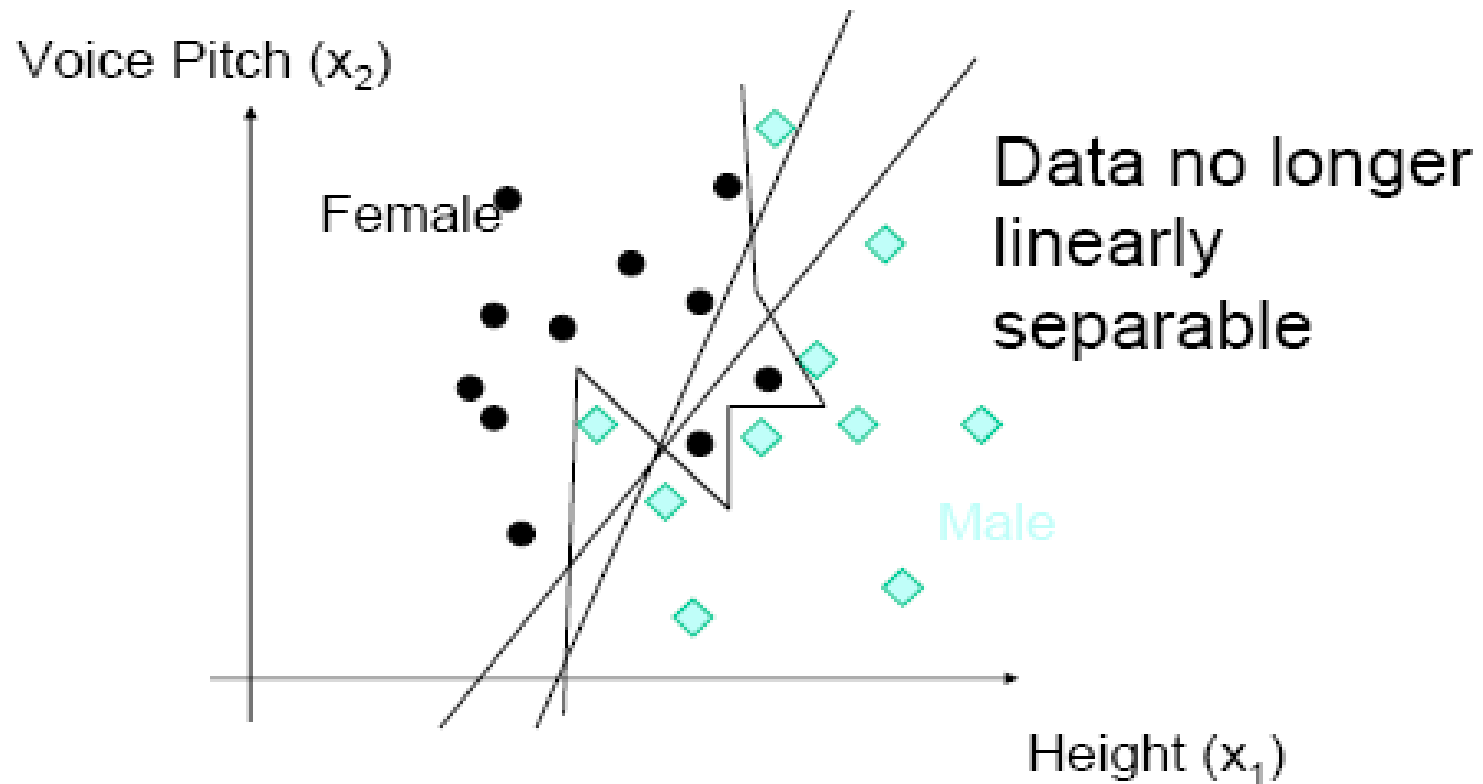
x_1	x_2	o_1	o_2	y
0	0	0	0	0
1	0	0	1	1
0	1	0	1	1
1	1	1	1	0

Piecewise linear classification using an MLP with threshold (perceptron) units

Multilayer Networks

- ▶ The single perceptron can only express linear decision surfaces.
- ▶ The kind of **multilayer networks** learned by the **back propagation** algorithm are capable of expressing a rich variety of nonlinear decision surfaces.

Multilayer Networks... Example



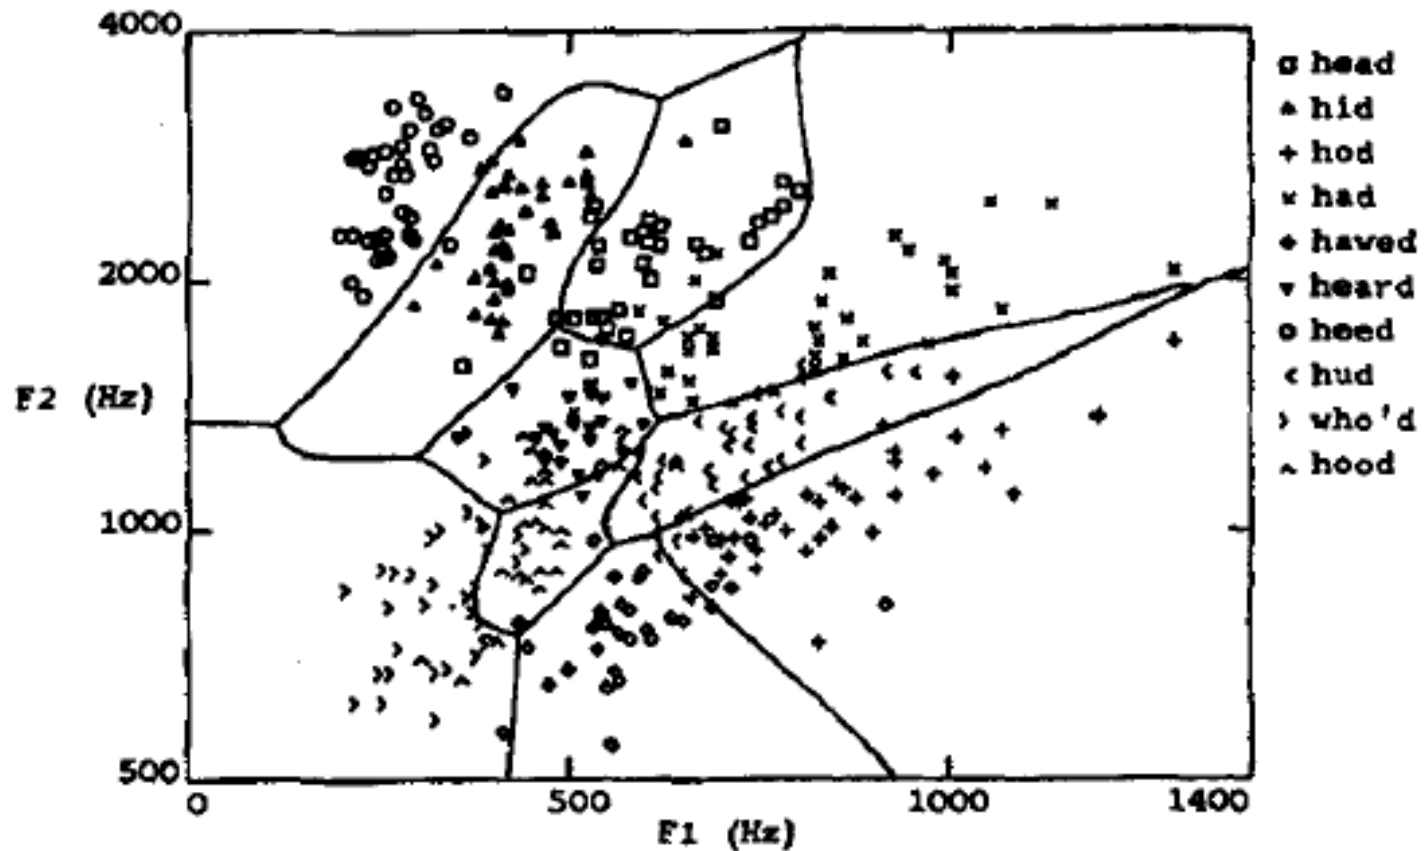
What is a good decision boundary ?

Multilayer Networks... Example

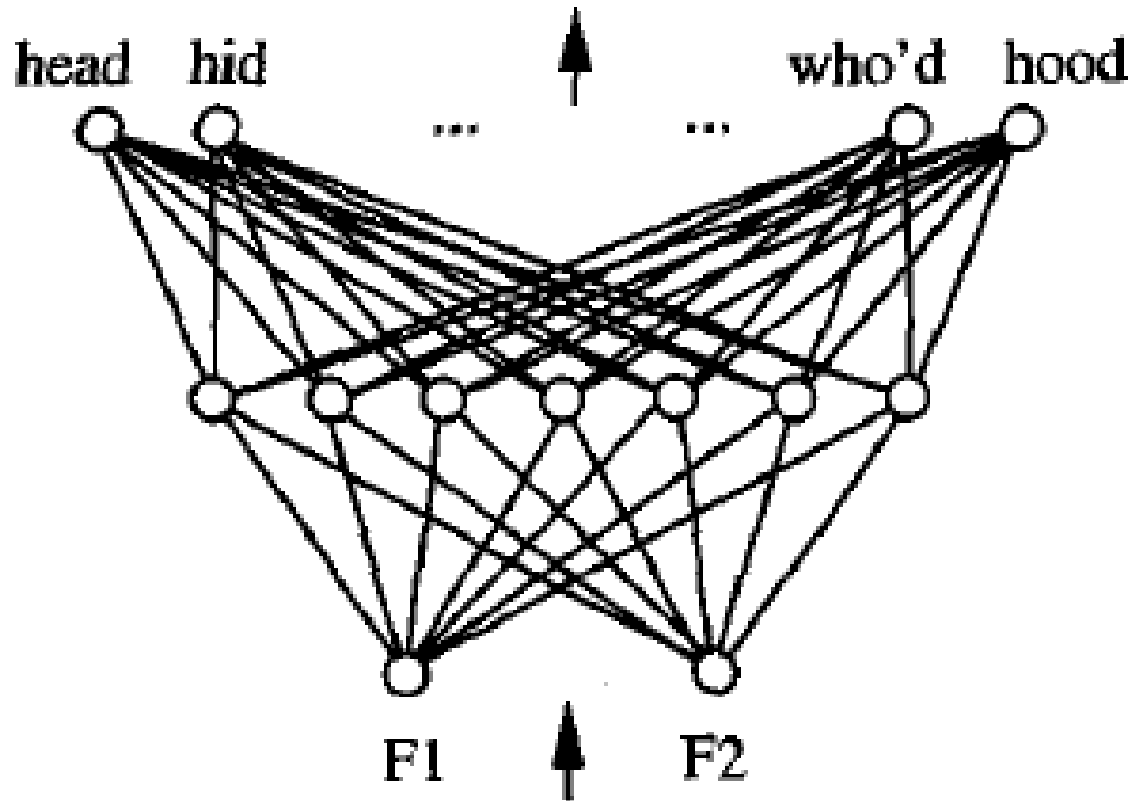
Example:

- ▶ The speech recognition task involves distinguishing among 10 possible vowels, all spoken in the context of "h-d" (i.e., "hid," "had," "head," "hood," etc.).

Multilayer Networks... Example



Multilayer Networks... Example



Reading Material

- ▶ **Artificial Intelligence, A Modern Approach**
Stuart J. Russell and Peter Norvig
 - Chapter 18.
- ▶ **Machine Learning**
Tom M. Mitchell
 - Chapter 4.

